

Using either your HTTP client or curl and your original HTTP server from Assignment 1, do the following:

- Place four different large files on the server. This can be done simply by copying the files to the server's directory, and then running the server. The files should be around 4 MiB long.
  - Start httpserver.
  - Start four separate instances of the client at the same time, one GETting each of the files and measure (using `time(1)`) how long it takes to get the files. Perhaps the best way to do this is to write a simple shell script (command file) that starts four copies of the client program in the background, by using `&` at the end.
- 
- Repeat the same experiment after you implement multi-threading. Is there any difference in performance?
- 
- What is likely to be the bottleneck in your system? How much concurrency is available in various parts, such as dispatch, worker, logging? Can you increase concurrency in any of these areas and, if so, how?